

Assignment 4

Fall 2016

CS834 Introduction to Information Retrieval

Dr. Michael Nelson

Mathew Chaney

December 7, 2016

Contents

1	Question 8.3	3
1.1	Question	3
1.2	Approach	3
1.3	Results	4
2	Question 8.4	5
2.1	Question	5
2.2	Approach	5
2.3	Results	5
3	Question 8.5	8
3.1	Question	8
3.2	Approach	8
3.3	Results	8
4	Question 8.7	9
4.1	Question	9
4.2	Approach	9
4.3	Results	9
5	Question 8.9	10
5.1	Question	10
5.2	Approach	10
5.3	Results	11
6	Appendix	12
6.1	Code listings	12
6.2	Script Output	18
7	References	21

List of Figures

1	Uninterpolated Recall-Precision Graph for CACM Queries 6 and 8.	5
2	Graph of interpolated precision at standard recall values for CACM queries 6 and 8. .	6
3	Average interpolated recall-precision graph for CACM Queries 6 and 8.	7
4	Recall-precision graph for all CACM Queries.	8

List of Tables

1	Calculations for CACM query 10 from all retrieved documents.	4
2	Interpolated precision at standard recall values for CACM queries 6 and 8.	6
3	Calculations for all CACM queries from all retrieved documents.	8
4	Table of BPREF values for a selection of CACM queries.	11

1 Question 8.3

1.1 Question

For one query in the CACM collection (provided at the book website), generate a ranking using Galago, and then calculate average precision, NDCG at 5 and 10, precision at 10, and the reciprocal rank by hand.

1.2 Approach

Galago version 3.10 was first downloaded from the Project Lemur Source Forge website, which can be found at <https://sourceforge.net/projects/lemur/files/lemur/galago-3.10/>. The CACM document corpus was downloaded from the textbook website, found at <http://www.search-engines-book.com/collections/>. Galago was used to create an index of the CACM corpus and then as a search server to respond to queries on that index.

The `getrel.py` and `q83.py` scripts (found in Listings 5 and 6, respectively) were created to issue CACM queries to the Galago search server using the Python Requests library [1]. The HTML responses the server sent back were parsed using the Python BeautifulSoup library [2], where the CACM document identifiers were extracted for use in calculating the different evaluation scores for the ranking. Various functions were added to the `getrel.py` script to derive the scores for all of the exercises in this assignment.

CACM query 10 was used for this exercise and the retrieval set was limited to the first 100 documents.

1.2.1 Initial Precision and Recall Calculations

Precision and Recall were calculated with the following equations:

$$Recall = \frac{|A \cap B|}{|A|}$$
$$Precision = \frac{|A \cap B|}{|B|}$$

Where A is the relevant set of documents for the query, and B is the set of retrieved documents.

1.2.2 Calculating Precision at Specific Rankings

A list of precision values was created by calculating the cumulative precision at each document ranking with the set of retrieved documents up to that ranking.

1.2.3 Calculating Average Precision

Average precision was calculated by adding the precision at each retrieval ranking position for documents which are part of $A \cap B$, or the set of retrieved documents that are relevant, and then dividing by the size of that set to obtain the average. This can also be described as the area under the precision-recall curve, which can be expressed as the following summation:

$$AveP = \sum_{k=1}^n P(k) \Delta r(k)$$

where k is the rank in the sequence of retrieved documents, n is the number of retrieved documents, $P(k)$ is the precision at cut-off k in the list, and $\Delta r(k)$ is the change in recall from items $k - 1$ to k .

1.2.4 Calculating Normalized Discounted Cumulative Gain (NDCG)

First, discounted cumulative gain at rank p (DCG_p) was calculated with the following formula:

$$DCG_p = rel_1 + \sum_{i=2}^p \frac{rel_i}{\log_2 i}$$

with rel_i being the relevancy for document i in the retrieval ranking. For this experiment, this value is either 0 or 1.

The ideal discounted cumulative gain at rank p ($IDCG_p$) is expressed as:

$$IDCG_p = 1 + \sum_{i=2}^p \frac{1}{\log_2 i}$$

Finally, normalized discounted cumulative gain at rank p ($NDCG_p$) is expressed as:

$$NDCG_p = \frac{DCG_p}{IDCG_p}$$

1.2.5 Calculating Reciprocal Rank

Reciprocal rank is defined as the reciprocal of the rank at which the first relevant document is found, so if the 3rd document in the retrieval ranking list is the first relevant document, the reciprocal rank is $\frac{1}{3}$.

1.3 Results

After building the index, CACM query 10 was processed by the `getrel.py` script, the output of which can be found in Listing 1. This script calculates all the values shown in Table 1.

```
1 [mchaney@mchaney-l getrel]$ python q83.py -q 10 -n 10000
2 query 10
3 query: parallel languages languages for parallel computation
4 precision: 0.0190816935003
5 recall: 0.914285714286
6 precision @10: 0.9
7 NDCG @5: 1.0
8 NDCG @10: 0.942709999032
9 avg precision: 0.5922383982
10 reciprocal rank: 1.0
```

Listing 1: Output from running the q83.py script for query 10 from the CACM collection.

Query #	Avg. Prec.	NDCG @5	NDCG @10	Prec. @10	Recip. Rank
10	0.5922383982	1.0	0.942709999032	0.9	1.0

Table 1: Calculations for CACM query 10 from all retrieved documents.

2 Question 8.4

2.1 Question

For two queries in the CACM collection, generate two uninterpolated recall-precision graphs, a table of interpolated precision values at standard recall levels, and the average interpolated recall-precision graph.

2.2 Approach

The `getrel.py`, `q84.py` and `graphs.R` scripts, found in Listings 5, 7 and 11, were used to complete this task.

2.3 Results

2.3.1 Uninterpolated Recall-Precision Graph

The uninterpolated recall-precision graph is shown in Figure 1.




Figure 1: Uninterpolated Recall-Precision Graph for CACM Queries 6 and 8.

2.3.2 Interpolated Precision

The graph for the interpolated precision at standard recall values is shown in Figure 2 and the table of the values for each query, including the averages, is shown in Table 2.




Figure 2: Graph of interpolated precision at standard recall values for CACM queries 6 and 8.

Recall	0.0	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1.0
Query 6	0.333	0.333	0.333	0.333	0.333	0.333	0.333	0.333	0.333	0.333	0.333
Query 8	0.5	0.5	0.5	0.5	0.4	0.4	0.4	0.4	0.4	0.4	0.4
Average	0.417	0.417	0.417	0.417	0.367	0.367	0.367	0.367	0.367	0.367	0.367

Table 2: Interpolated precision at standard recall values for CACM queries 6 and 8.

2.3.3 Average Interpolated Precision

The graph of the average interpolated precision at standard recall values for CACM queries 6 and 8 can be found in Figure 3.




Figure 3: Average interpolated recall-precision graph for CACM Queries 6 and 8.

3 Question 8.5

3.1 Question

Generate the mean average precision, recall-precision graph, average NDCG at 5 and 10, and precision at 10 for the entire CACM query set.

3.2 Approach

The `getrel.py` and `q85.py` scripts, found in Listings 5 and 8, were used to complete this question.

3.3 Results

Using only queries for which relevance judgments exist the MAP, NDCG at 5 and at 10, and the precision at 10 were calculated over all retrieved documents. The results can be found in Table 3. The generated recall-precision graph for the entire query set can be found in Figure 4.

MAP	NDCG @5	NDCG @10	Prec. @10
0.339552098123	0.461648777763	0.381724764912	0.317647058824

Table 3: Calculations for all CACM queries from all retrieved documents.




Figure 4: Recall-precision graph for all CACM Queries.

4 Question 8.7

4.1 Question

Another measure that has been used in a number of evaluations is R-precision. This is defined as the precision at R documents, where R is the number of relevant documents for a query. It is used in situations where there is a large variation in the number of relevant documents per query. Calculate the average *R-precision* for the CACM query set and compare it to the other measures.

4.2 Approach

The `getrel.py` and `q87.py` scripts were used to complete this exercise. They can be found in Listings 5 and 9. The script was run over the entire document set, which will minimize precision and maximize recall. This will be taken into consideration when comparing the R-precision score to other measures.

4.3 Results

The output of running the `q87.py` script can be found in Listing 14.

```
1 [mchaney@mchaney-l getrel]$ python q87.py -n 3204 -q 10
2 query 10
3 query: parallel languages languages for parallel computation
4 relevant: 35
5 retrieved: 1677
6 precision: 0.0190816935003
7 recall: 0.914285714286
8 precision @10: 0.9
9 NDCG @5: 1.0
10 NDCG @10: 0.942709999032
11 avg precision: 0.5922383982
12 reciprocal rank: 1.0
13 |R|: 35
14 R-precision: 0.555555555556
```

Listing 2: Output of `q87.py` for query 10.

4.3.1 Comparison

For query 10 the R-precision score was 0.5, which is very close to the average precision over the entire document set. This rather simplistic comparison is one mark towards R-precision being a semi-reliable measure of performance for a ranking. R-precision seems to be somewhat of an intuitive normalization of precision, which the textbook touches on. This measure will favor rankings that push more relevant documents into higher ranks, which sounds like the makings of a strong effectiveness measure, but it is not without faults.

The only problem with this measure is if the relevant set is very small the results of the calculation may vary wildly. This may make the measure inappropriate for a targeted search, because it will tend toward either 0 or 1, which isn't at all useful for comparing rankings.

5 Question 8.9

5.1 Question

For one query in the CACM collection, generate a ranking and calculate BPREF. Show that the two formulations of BPREF give the same value.

5.2 Approach

The first version of BPREF found in the text book is defined as follows:

$$BPREF = \frac{1}{R} \sum_{d_r} (1 - \frac{N_{d_r}}{R}) \quad (1)$$

Where d_r is a relevant document, N_{d_r} gives the number of non-relevant documents that are ranked higher than document d_r and R is the size of the set of all relevant documents for the given query.

The second form is:

$$BPREF = \frac{P}{P + Q} \quad (2)$$

Where P is the number of preferences that agree and Q is the number of preferences that disagree.

The scripts `getrel.py` and `q89.py`, found in Listings 5 and 10 were used to complete this exercise.

5.2.1 Implementation

The script `q89.py` was used as a driver program for this exercise and can be found in Listing 10.

The first BPREF equation (1) was implemented as the `bpref1` function, which can be found in Listing 3. This function iterates over the given set of relevant documents and calculates the number of non-relevant documents

```
129 def bpref1(rel, retr):
130     """calculate BPREF as the first equation in the textbook"""
131     relset = set(rel)
132     retrset = set(retr)
133     R = float(len(rel))
134     res = 0
135     for dr in rel:
136         if dr not in retr:
137             Ndr = R
138         else:
139             Ndr = float(len([doc for doc in retr[:retr.index(dr)] if doc not in rel]))
140             res += (1.0 - Ndr/R)
141     return 1.0/R * res
```

Listing 3: bpref1 function

The second BPREF equation (1) was implemented as the `bpref2` function, and can be found in Listing 3.

```

143 def bpref2(rel, retr):
144     """calculate BPREF as the second equation in the textbook"""
145     p = 0.0
146     q = 0.0
147     for dr in rel:
148         if dr not in retr:
149             p += 0
150             q += len(retr)
151         else:
152             p += float(len([doc for doc in retr[retr.index(dr)]: if doc not in rel]))
153             q += float(len([doc for doc in retr[:retr.index(dr)]: if doc not in rel]))
154     return p / (p + q)

```

Listing 4: `bpref2` function

5.3 Results

Considering one example did not seem sufficient, a number of other queries were tested with the results shown in Table 4.

Query	R	BPREF(1)	BPREF(2)
1	5	0.2	0.172413793103
5	8	0.078125	0.0704225352113
8	3	0.222222222222	0.181818181818
11	19	0.443213296399	0.354767184035
16	17	0.0692041522491	0.0626959247649
21	11	0.380165289256	0.304635761589
37	12	0.0833333333333	0.0774193548387
42	21	0.106575963719	0.0902111324376

Table 4: Table of BPREF values for a selection of CACM queries.

6 Appendix

6.1 Code listings

```
1 import argparse
2 import re
3 import requests
4 import sys
5 import xmltodict
6 import numpy as np
7 from math import log
8 from bs4 import BeautifulSoup
9
10 def parseargs():
11     """parse command line arguments"""
12     parser = argparse.ArgumentParser()
13     parser.add_argument('-p', '--port', type=int, default=42247, help='galago server port')
14     parser.add_argument('-q', '--qnum', nargs='+', type=int, default=[6, 8], help='query
15         number')
16     parser.add_argument('-n', type=int, default=10, help='number of retrieval pages')
17     return parser.parse_args()
18
19 args = parseargs()
20
21 def buildrel():
22     """read CACM relevance judgments for each query"""
23     rel = {}
24     for line in open('cacm.rel').readlines():
25         q, _, doc, _ = line.split()
26         if q not in rel:
27             rel[q] = []
28         rel[q].append(int(doc.split('-')[1]))
29     return rel
30
31 def buildqueries():
32     """read CACM query file into a map structure for later use"""
33     with open('cacm.query.xml') as fd:
34         return xmltodict.parse(fd.read())
35
36 REL = buildrel()
37 QUERIES = buildqueries()
38 RE = re.compile('/home/mchaney/workspace/edu/cs834-f16/assignments/assignment4/code/cacm/
39 docs/CACM-([d]+).html')
40 ID = {'id': 'result'}
41 URL = 'http://0.0.0.0:{0}/search'
42 QUERY1 = 'what articles exist which deal with tss time sharing system an operating system
43 for ibm computers'
44 PDICT = {'q': QUERY1, 'start': 0, 'n': args.n}
45
46 def query(qstr, port=args.port):
47     """issue a query to a running Galago server"""
48     PDICT['q'] = qstr
49     PDICT['n'] = args.n
50     res = requests.get(URL.format(port), params=PDICT)
51     if not res.ok:
52         return None
53     soup = BeautifulSoup(res.text, 'html.parser')
54     return [int(RE.match(href.text).groups()[0]) for href in soup.select("#result a")]
55
56 def recall(rel, retr):
57     """calculate the recall given relevant and retrieved documents"""
58     relset = set(rel)
59     retrset = set(retr)
60     return float(len(relset.intersection(retrset))) / len(relset)
61
62 def precision(rel, retr):
63     """calculate the precision given relevant and retrieved documents"""
64     relset = set(rel)
65     retrset = set(retr)
66     return float(len(relset.intersection(retrset))) / len(retrset)
67
68 def run(rel, retr, func):
69     """calculate recall or precision at all ranks given relevant and retrieved documents"""
70     rr = []
71     for i in range(1, len(retr)+1):
```

```

69         rr.append(func(rel, retr[:i]))
70     return rr
71
72 def avg(rel, retr, func):
73     """calculate avg recall or precision given relevant and retrieved documents"""
74     prun = run(rel, retr, func)
75     res = []
76     for i in range(len(retr)):
77         if retr[i] in rel:
78             res.append(prun[i])
79     if len(res) == 0:
80         return 0.0
81     return float(sum(res))/len(res)
82
83 def getrel(rel, retr, i):
84     """get a relevancy score for a given retrieved document"""
85     return 1 if retr[i] in rel else 0
86
87 def DCG(rel, retr, p):
88     """calculate DCG at p given relevant and retrieved documents"""
89     sum = 0
90     for i in range(2, p+1):
91         sum += float(getrel(rel, retr, i-1)) / log(i, 2)
92     return getrel(rel, retr, 0) + sum
93
94 def IDCG(p):
95     """calculate IDCG at p given relevant and retrieved documents"""
96     sum = 0
97     for i in range(2, p+1):
98         sum += 1 / log(i, 2)
99     return 1 + sum
100
101 def NDCG(rel, retr, p):
102     """calculate NDCG at p given relevant and retrieved documents"""
103     dcg = DCG(rel, retr, p)
104     idcg = IDCG(p)
105     return dcg / idcg
106
107 def reciprank(rel, retr):
108     """calculate reciprocal rank given relevant and retrieved documents"""
109     for i in range(1, len(retr)+1):
110         if retr[i-1] in rel:
111             return 1.0 / i
112     return 0.0
113
114 def ipr(rrun, prun):
115     """calculate interpolated precision at std recall given recall and precision runs"""
116     res = []
117     for i in np.arange(0, 1.1, .1):
118         for j in range(len(rrun)):
119             if rrun[j] > i:
120                 idx = j
121                 break
122         res.append(max(prun[idx:]))
123     return np.arange(0, 1.1, 0.1), res
124
125 def rprecision(rel, prun):
126     """calculate R-precision"""
127     return prun[len(rel)]
128
129 def bpref1(rel, retr):
130     """calculate BPREF as the first equation in the textbook"""
131     relset = set(rel)
132     retrset = set(retr)
133     R = float(len(rel))
134     res = 0
135     for dr in rel:
136         if dr not in retr:
137             Ndr = R
138         else:
139             Ndr = float(len([doc for doc in retr[:retr.index(dr)] if doc not in rel]))
140         res += (1.0 - Ndr/R)
141     return 1.0/R * res
142
143 def bpref2(rel, retr):
144     """calculate BPREF as the second equation in the textbook"""
145     p = 0.0

```

```

146 q = 0.0
147 for dr in rel:
148     if dr not in retr:
149         p += 0
150         q += len(retr)
151     else:
152         p += float(len([doc for doc in retr[retr.index(dr):] if doc not in rel]))
153         q += float(len([doc for doc in retr[:retr.index(dr)] if doc not in rel]))
154     return p / (p + q)
155
156 def getquery(qnum):
157     """get the query string"""
158     return QUERIES['parameters'][ 'query' ][qnum-1][ 'text' ]
159
160 def process(qnum):
161     """run a number of calculations used for some exercises"""
162     qstr = getquery(qnum)
163     retr = query(qstr)
164     if str(qnum) not in REL:
165         return [None]*12
166     rel = REL[str(qnum)]
167     prun = run(rel, retr, precision)
168     rrun = run(rel, retr, recall)
169     prec = precision(rel, retr)
170     rec = recall(rel, retr)
171     avgprec = avg(rel, retr, precision)
172     ndcg5 = NDCG(rel, retr, 5)
173     ndcg10 = NDCG(rel, retr, 10)
174     recip = reciprank(rel, retr)
175     return qnum, qstr, retr, rel, prun, rrun, prec, rec, ndcg5, ndcg10, avgprec, recip
176
177 def printresults(qnum, qstr, retr, rel, prun, rrun, prec, rec, ndcg5, ndcg10, avgprec, recip):
178     if not qnum:
179         return
180     print 'query {0}'.format(qnum)
181     print 'query: {0}'.format(qstr)
182     if args.n == 10:
183         print 'relevant: {0}'.format(len(rel))
184         print 'retrieved: {0}'.format(len(retr))
185         print 'p-run: {0}'.format(prun)
186         print 'r-run: {0}'.format(rrun)
187     else:
188         print 'relevant: {}'.format(len(rel))
189         print 'retrieved: {}'.format(len(retr))
190     print 'precision: {0}'.format(prec)
191     print 'recall: {0}'.format(rec)
192     print 'precision @10: {0}'.format(prun[9])
193     print 'NDCG @5: {0}'.format(ndcg5)
194     print 'NDCG @10: {0}'.format(ndcg10)
195     print 'avg precision: {0}'.format(avgprec)
196     print 'reciprocal rank: {0}'.format(recip)
197
198 def printdata(rrun, prun, fname):
199     """print a data table for Rscript consumption"""
200     with open(fname, 'w') as fd:
201         zipped = zip(rrun, prun)
202         for z in zipped:
203             fd.write('{0}\t{1}\n'.format(z[0], z[1]))
204
205 if __name__ == '__main__':
206     for qnum in args.qnum:
207         printresults(*process(qnum))

```

Listing 5: getrel.py

```

1 from getrel import *
2
3 TABLE = """\begin{table}[h!]
4 \centering
5 \begin{tabular}{|c|c|c|c|c|c|c|}
6 \hline
7 Query \# & Avg. Prec. & NDCG @5 & NDCG @10 & Prec. @10 & Recip. Rank & \\
8 \hline
9 {0} & {1} & {2} & {3} & {4} & {5} & \\
10 \hline

```

```

11 \\end{{tabular}}
12 \\caption{{Calculations for CACM query {6} from all retrieved documents.}}
13 \\label{{tab:q83}}
14 \\end{{table}}
15 """
16
17 def printtab(qnum, qstr, retr, rel, prun, rrun, prec, rec, ndcg5, ndcg10, avgprec, recip,
18             rprec):
19     fname = 'query{0}.tab'.format(qnum)
20     with open(fname, 'w') as fd:
21         fd.write(TABLE.format(qnum, avgprec, ndcg5, ndcg10, prun[9], recip, qnum))
22
23 for qnum in args.qnum:
24     results = process(qnum)
25     printresults(*results)
26     printtab(*results)

```

Listing 6: q83.py

```

1 from getrel import *
2
3 HEAD = """\\begin{table}[H]
4 \\centering
5 \\begin{tabular}{l l l l l l l l l l l l}
6 Recall & 0.0 & 0.1 & 0.2 & 0.3 & 0.4 & 0.5 & 0.6 & 0.7 & 0.8 & 0.9 & 1.0 \\
7 \\cline{2-12}
8 """
9
10 ROW = """{0} & {1:.3g} & {2:.3g} & {3:.3g} & {4:.3g} & {5:.3g} & {6:.3g} & {7:.3g} & {8:.3g}
11 & {9:.3g} & {10:.3g} & {11:.3g} \\
12 \\cline{2-12}
13 """
14
15 TAIL = """\\end{table}
16 \\caption{.}
17 \\label{tab:ipr68}
18 \\end{table}
19 """
20
21 def printtable(iprl):
22     with open('iptab.tex', 'w') as fd:
23         fd.write(HEAD)
24         for iprun, qnum in iprl:
25             fd.write(ROW.format('Query {0}'.format(qnum), *iprun))
26             avg = [float(sum(col))/len(col) for col in zip(*[col[0] for col in iprl])]
27             printdata(np.arange(0, 1.1, .1), avg, 'avg.dat')
28             fd.write(ROW.format('Average', *avg))
29             fd.write(TAIL)
30
31 iprl = []
32 for qnum in args.qnum:
33     results = process(qnum)
34     printresults(*results)
35     qnum, qstr, retr, rel, prun, rrun, prec, rec, ndcg5, ndcg10, avgprec, recip = results
36     printdata(rrun, prun, 'urpg{0}.dat'.format(qnum))
37     irr, iprun = ipr(rrun, prun)
38     printdata(irr, iprun, 'ipr{0}.dat'.format(qnum))
39     iprl.append((iprun, qnum))
40
41 printtable(iprl)

```

Listing 7: q84.py

```

1 from getrel import *
2
3 TABLE = """\\begin{{table}}[h!]
4 \\centering
5 \\begin{{tabular}}{{l | c | c | c | c | }}
6 \\hline
7 MAP & NDCG @5 & NDCG @10 & Prec. @10 \\
8 \\hline
9 {0} & {1} & {2} & {3} \\
10 \\hline
11 \\end{{tabular}}
12 \\caption{{Calculations for all CACM queries from all retrieved documents.}}

```

```

13 \\label{{tab:q85}}
14 \\end{{table}}
15 """
16
17 def printtab(fname, cacmmmap, avgndcg5, avgndcg10, avgprec10):
18     with open(fname, 'w') as fd:
19         fd.write(TABLE.format(cacmmmap, avgndcg5, avgndcg10, avgprec10))
20
21 netavg = []
22 iprl = []
23 ndcg5lst = []
24 ndcg10lst = []
25 prunlst = []
26 for i in range(1, 64):
27     qnum, qstr, retr, rel, prun, rrun, prec, rec, ndcg5, ndcg10, avgprec, recip = process(i)
28     if avgprec:
29         netavg.append(avgprec)
30         prunlst.append(prun[9])
31         irrun, iprun = ipr(rrun, prun)
32         iprl.append((iprun, qnum))
33         ndcg5lst.append(ndcg5)
34         ndcg10lst.append(ndcg10)
35
36 # MAP
37 cacmmmap = float(sum(netavg)) / len(netavg)
38 print 'average precision: {0}'.format(cacmmmap)
39
40 # Recall-Precision
41 netavgrpg = [float(sum(col))/len(col) for col in zip(*[col[0] for col in iprl])]
42 printdata(np.arange(0, 1.1, .1), netavgrpg, 'avgq85.dat')
43
44 # NDCG @ 5 and 10
45 avgndcg5 = float(sum(ndcg5lst))/len(ndcg5lst)
46 avgndcg10 = float(sum(ndcg10lst))/len(ndcg10lst)
47 print 'NDCG @5 : {0}'.format(avgndcg5)
48 print 'NDCG @10: {0}'.format(avgndcg10)
49
50 # precision at 10
51 avgprec10 = float(sum(prunlst))/len(prunlst)
52 print 'Precision @10: {0}'.format(avgprec10)
53
54 printtab('q85.tab', cacmmmap, avgndcg5, avgndcg10, avgprec10)

```

Listing 8: q85.py

```

1 from getrel import *
2
3 for qnum in args.qnum:
4     qnum, qstr, retr, rel, prun, rrun, prec, rec, ndcg5, ndcg10, avgprec, recip = process(
5         qnum)
6     printresults(qnum, qstr, retr, rel, prun, rrun, prec, rec, ndcg5, ndcg10, avgprec, recip)
7     rprec = rprecision(rel, prun)
8     print '|R|: {}'.format(len(rel))
9     print 'R-precision: {}'.format(rprec)

```

Listing 9: q87.py

```

1 from getrel import *
2
3 HEAD = """\\begin{table}[H]
4 \\centering
5 \\begin{tabular}{| l | l | l | l |}
6 \\hline
7 Query & R & BPREF(1) & BPREF(2) \\\\
8 \\hline
9 """
10
11 ROW = """{ } & { } & { } & { }\\
12 \\hline
13 """
14
15 TAIL = """\\end{tabular}
16 \\caption{Table of BPREF values for a selection of CACM queries.}

```



```

17 \\label{tab:bpref}
18 \\end{table}
19 ""
20
21 def printtab(values):
22     with open('q89.tab', 'w') as fd:
23         fd.write(HEAD)
24         for qnum, R, b1, b2 in values:
25             fd.write(ROW.format(qnum, R, b1, b2))
26         fd.write(TAIL)
27
28 def getsub(rel, retr):
29     """bound the size of retrieved list by R non-relevant documents"""
30     sub = []
31     count = 0
32     for doc in retr:
33         if doc not in rel:
34             count += 1
35             if count > len(rel):
36                 break
37             sub.append(doc)
38     return sub
39
40 values = []
41 for qnum in args.qnum:
42     results = process(qnum)
43     retr, rel = results[2], results[3]
44     if not retr:
45         print 'no results for query {}'.format(qnum)
46         continue
47     sub = getsub(rel, retr)
48     b1 = bpref1(rel, sub)
49     b2 = bpref2(rel, sub)
50     print 'Query: {}'.format(qnum)
51     print 'R: {}'.format(len(rel))
52     print 'BPREF1: {}'.format(b1)
53     print 'BPREF2: {}'.format(b2)
54     values.append((qnum, len(rel), b1, b2))
55 printtab(values)

```

Listing 10: q89.py

```

1 plotone <- function(data, fname) {
2     pdf(fname)
3     plot(data, type='o', pch=15, ylim=c(0,1), xlim=c(0,1),
4           ylab="Precision", xlab="Recall")
5     dev.off()
6 }
7 urpgraph <- function(d1, d2, fname) {
8     pdf(fname)
9     plot(d1, lwd=2, type='o', pch=18, ylim=c(0,1), xlim=c(0,1), col="gray60",
10          ylab="Precision", xlab="Recall")
11     lines(d2, lwd=2, type="o", pch=15, col="gray30")
12     legend(0.8, 1, c('Query 6', 'Query 8'), cex=0.8,
13            col=c('gray60', 'gray30'), lty=c(1,1), pch=c(18,15))
14     dev.off()
15 }
16 iprgraph <- function(d1, d2, id1, id2, fname) {
17     pdf(fname)
18     plot(d1, lwd=2, type="p", pch=18, ylim=c(0,1), xlim=c(0,1), col="gray60",
19          ylab="Precision", xlab="Recall")
20     lines(id1, lwd=2, type="s", col="gray60")
21     lines(d2, lwd=2, type="p", pch=15, col="gray30")
22     lines(id2, lwd=2, type="s", col="gray30")
23     legend(0.8, 1, c('Query 6', 'Query 8'), cex=0.8,
24            col=c('gray60', 'gray30'), lty=c(1,1), pch=c(18,15))
25     dev.off()
26 }
27 aipgraph <- function(avg, id1, id2, fname) {
28     pdf(fname)
29     plot(avg, lwd=2, type="l", ylim=c(0,1), xlim=c(0,1), col="black",
30          ylab="Precision", xlab="Recall")
31     lines(avg, lwd=2, type="p", pch=17, col="black")
32     lines(id1, lwd=2, type="s", col="gray60")
33     lines(id2, lwd=2, type="s", col="gray30")
34     legend(0.8, 1, c('Query 6', 'Query 8', 'Average'), cex=0.8,

```

```

35     col=c('gray60', 'gray30', 'black'), lty=c(1,1,1), pch=c(18,15,17))
36     dev.off()
37 }
38
39 args = commandArgs(trailingOnly=TRUE)
40
41 d1 <- read.table(paste('urpg', args[1], '.dat', sep=''))
42 d2 <- read.table(paste('urpg', args[2], '.dat', sep=''))
43
44 plotone(d1, paste('urpg', args[1], '.pdf', sep=''))
45 plotone(d2, paste('urpg', args[2], '.pdf', sep=''))
46 urpgraph(d1, d2, paste('urpg', args[1], '', args[2], '.pdf', sep=''))
47
48 id1 <- read.table(paste('ipr', args[1], '.dat', sep=''))
49 id2 <- read.table(paste('ipr', args[2], '.dat', sep=''))
50 iprgraph(d1, d2, id1, id2, paste('ipr', args[1], '', args[2], '.pdf', sep=''))
51
52 avg <- read.table('avg.dat')
53 aipgraph(avg, id1, id2, paste('aipr', args[1], args[2], '.pdf', sep=''))
54
55 overallavg <- read.table('avgq85.dat')
56 plotone(overallavg, 'avgq85.pdf')

```

Listing 11: Script used to generate the recall-precision graphs

6.2 Script Output

```

1 [mchaney@mchaney-l getrel]$ python q84.py -q 6 8
2 query 6
3 query: interested in articles on robotics motion planning particularly the geometric and
4       combinatorial aspects we are not interested in the dynamics of arm motion
5 relevant: [1543, 2078, 2828]
6 retrieved: [2826, 695, 531, 2828, 2931, 1014, 2721, 2078, 1543, 1867]
7 p-run: [0.0, 0.0, 0.0, 0.25, 0.2, 0.16666666666666666, 0.14285714285714285, 0.25,
8         0.3333333333333333, 0.3]
9 r-run: [0.0, 0.0, 0.0, 0.3333333333333333, 0.3333333333333333, 0.3333333333333333,
10         0.3333333333333333, 0.6666666666666666, 1.0, 1.0]
11 precision: 0.3
12 recall: 1.0
13 precision @10: 0.3
14 NDCG @5: 0.140386094433
15 NDCG @10: 0.218631536782
16 avg precision: 0.2777777777777778
17 reciprocal rank: 0.25
18 query 8
19 query: addressing schemes for resources in networks resource addressing in network
20       operating systems
21 relevant: [2625, 2849, 3032]
22 retrieved: [2951, 2625, 2949, 2500, 3032, 2776, 2371, 1685, 1752, 2541]
23 p-run: [0.0, 0.5, 0.3333333333333333, 0.25, 0.4, 0.3333333333333333, 0.2857142857142857,
24         0.25, 0.2222222222222222, 0.2]
25 r-run: [0.0, 0.3333333333333333, 0.3333333333333333, 0.3333333333333333, 0.6666666666666666,
26         0.6666666666666666, 0.6666666666666666, 0.6666666666666666, 0.6666666666666666,
27         0.6666666666666666]
28 precision: 0.2
29 recall: 0.6666666666666667
30 precision @10: 0.2
31 NDCG @5: 0.40169418877
32 NDCG @10: 0.272276725167
33 avg precision: 0.45
34 reciprocal rank: 0.5

```

Listing 12: Output of the q84.py script.

```

1 [mchaney@mchaney-l getrel]$ python q85.py -n 3204
2 average precision: 0.339552098123
3 NDCG @5 : 0.461648777763
4 NDCG @10: 0.381724764912
5 Precision @10: 0.317647058824

```

Listing 13: Output of the q85.py script.

```

1 [mchaney@mchaney-l getrel]$ python q87.py -n 3204 -q 10
2 query 10
3 query: parallel languages languages for parallel computation
4 relevant: 35
5 retrieved: 1677
6 precision: 0.0190816935003
7 recall: 0.914285714286
8 precision @10: 0.9
9 NDCG @5: 1.0
10 NDCG @10: 0.942709999032
11 avg precision: 0.5922383982
12 reciprocal rank: 1.0
13 |R|: 35
14 R-precision: 0.555555555556

```

Listing 14: Output of the q87.py script.

```

1 [mchaney@mchaney-l getrel]$ python q89.py -n 11 -q 1 5 6 8 10 16 21
2 query 1
3 query: what articles exist which deal with tss time sharing system an operating system
4 for ibm computers
5 relevant: 5
6 retrieved: 11
7 precision: 0.0909090909091
8 recall: 0.2
9 precision @10: 0.1
10 NDCG @5: 0.280772188866
11 NDCG @10: 0.190313263771
12 avg precision: 1.0
13 reciprocal rank: 1.0
14 BPREF1: 0.2
15 BPREF2: 0.172413793103
16 query 5
17 query: id like papers on design and implementation of editing interfaces window managers
18 command interpreters etc the essential issues are human interface design with views
19 on improvements to user efficiency effectiveness and satisfaction
20 relevant: 8
21 retrieved: 11
22 precision: 0.181818181818
23 recall: 0.25
24 precision @10: 0.1
25 NDCG @5: 0.140386094433
26 NDCG @10: 0.0951566318853
27 avg precision: 0.215909090909
28 reciprocal rank: 0.25
29 BPREF1: 0.078125
30 BPREF2: 0.0704225352113
31 query 6
32 query: interested in articles on robotics motion planning particularly the geometric and
33 combinatorial aspects we are not interested in the dynamics of arm motion
34 relevant: 3
35 retrieved: 11
36 precision: 0.272727272727
37 recall: 1.0
38 precision @10: 0.3
39 NDCG @5: 0.140386094433
40 NDCG @10: 0.218631536782
41 avg precision: 0.277777777778
42 reciprocal rank: 0.25
43 BPREF1: 0.0
44 BPREF2: 0.0
45 query 8
46 query: addressing schemes for resources in networks resource addressing in network
47 operating systems
48 relevant: 3
49 retrieved: 11
50 precision: 0.181818181818
51 recall: 0.666666666667
52 precision @10: 0.2
53 NDCG @5: 0.40169418877
54 NDCG @10: 0.272276725167
55 avg precision: 0.45
56 reciprocal rank: 0.5
57 BPREF1: 0.222222222222
58 BPREF2: 0.181818181818

```

```

54 query 10
55 query: parallel languages languages for parallel computation
56 relevant: 35
57 retrieved: 11
58 precision: 0.818181818182
59 recall: 0.257142857143
60 precision @10: 0.9
61 NDCG @5: 1.0
62 NDCG @10: 0.942709999032
63 avg precision: 1.0
64 reciprocal rank: 1.0
65 BPREF1: 0.257142857143
66 BPREF2: 0.0592105263158
67 query 16
68 query: find all descriptions of file handling in operating systems based on multiple
        processes and message passing
69 relevant: 17
70 retrieved: 11
71 precision: 0.181818181818
72 recall: 0.117647058824
73 precision @10: 0.1
74 NDCG @5: 0.0
75 NDCG @10: 0.0736232203436
76 avg precision: 0.174242424242
77 reciprocal rank: 0.166666666667
78 BPREF1: 0.0692041522491
79 BPREF2: 0.0218579234973
80 query 21
81 query: computational complexity intractability class complete reductions algorithms and
        efficiency
82 relevant: 11
83 retrieved: 11
84 precision: 0.454545454545
85 recall: 0.454545454545
86 precision @10: 0.5
87 NDCG @5: 0.682466377636
88 NDCG @10: 0.590418091601
89 avg precision: 0.745396825397
90 reciprocal rank: 1.0
91 BPREF1: 0.380165289256
92 BPREF2: 0.21875

```

Listing 15: Output of the q89.py script.

7 References

- [1] Kenneth Reitz. Requests: HTTP for Humans. Available at <http://docs.python-requests.org/en/master/>. Accessed: 2016/09/20.
- [2] Leonard Richardson. Beautiful Soup. Available at: <https://www.crummy.com/software/beautifulsoup/>. Accessed: 2016/09/20.