

Assignment 4

Fall 2016

CS834 Introduction to Information Retrieval

Dr. Michael Nelson

Mathew Chaney

December 5, 2016

Contents

1	Question 8.3	3
1.1	Question	3
1.2	Approach	3
1.3	Results	4
2	Question 8.4	5
2.1	Question	5
2.2	Approach	5
2.3	Results	5
3	Question 8.5	8
3.1	Question	8
3.2	Approach	8
3.3	Results	8
4	Question 8.7	9
4.1	Question	9
4.2	Approach	9
4.3	Results	9
5	Appendix	10
5.1	Code listings	10
6	References	16

Listings

1	Output from running the getrel.py script for queries 1 and 10 from the CACM collection.	4
2	Output of q87.py for query 10.	9
3	getrel.py	10
4	q83.py	12
5	q84.py	12
6	q85.py	13
7	q87.py	14
8	Script used to generate the recall-precision graphs	14

List of Figures

1	Uninterpolated Recall-Precision Graph for CACM Queries 6 and 8.	5
2	Graph of interpolated precision at standard recall values for CACM queries 6 and 8.	6
3	Average interpolated recall-precision graph for CACM Queries 6 and 8.	7
4	Recall-precision graph for all CACM Queries.	8

List of Tables

1	Calculations for CACM query 10 from all retrieved documents.	4
2	Interpolated precision at standard recall values for CACM queries 6 and 8.	6
3	Calculations for all CACM queries from all retrieved documents.	8

1 Question 8.3

1.1 Question

For one query in the CACM collection (provided at the book website), generate a ranking using Galago, and then calculate average precision, NDCG at 5 and 10, precision at 10, and the reciprocal rank by hand.

1.2 Approach

Galago version 3.10 was first downloaded from the Project Lemur Source Forge website, which can be found at the following URL: <https://sourceforge.net/projects/lemur/files/lemur/galago-3.10/>. The CACM document corpus was downloaded from the textbook's website, found here: <http://www.search-engines-book.com/collections/>. Galago was used to create an index of the CACM corpus and to run as a server to respond to queries on that index.

The `getrel.py` and `q83.py` scripts (found in Listings 3 and 4, respectively) was created to issue queries to the Galago search server using the Python Requests library [1]. The HTML responses were then parsed using the Python BeautifulSoup library [2], where the CACM document identifiers were extracted for use in calculating the different evaluation scores for the Galago ranking.

The query used was from the CACM query set, number 10, and only the first 1000 retrieved documents were considered when calculating all scores for this experiment.

1.2.1 Initial Precision and Recall Calculations

Precision and Recall were calculated with the following equations:

$$Recall = \frac{|A \cap B|}{|A|}$$
$$Precision = \frac{|A \cap B|}{|B|}$$

In these equations, A is the relevant set of documents for the query, and B is the set of retrieved documents.

1.2.2 Calculating Precision at Specific Rankings

A list of precision values was created by calculating the cumulative precision at each document ranking with the set of retrieved documents up to that ranking.

1.2.3 Calculating Average Precision

Average precision was calculated by adding the precision at each retrieval ranking position for documents which are part of $A \cap B$, or the set of retrieved documents that are relevant, and then dividing by the size of that set to obtain the average. This can also be described as the area under the precision-recall curve, which can be expressed as the following summation:

$$AveP = \sum_{k=1}^n P(k) \Delta r(k)$$

where k is the rank in the sequence of retrieved documents, n is the number of retrieved documents, $P(k)$ is the precision at cut-off k in the list, and $\Delta r(k)$ is the change in recall from items $k - 1$ to k .

1.2.4 Calculating Normalized Discounted Cumulative Gain (NDCG)

First, discounted cumulative gain at rank p (DCG_p) was calculated with the following formula:

$$DCG_p = rel_1 + \sum_{i=2}^p \frac{rel_i}{\log_2 i}$$

The ideal discounted cumulative gain at rank p ($IDCG_p$) is a simple series, expressed as:

$$IDCG_p = 1 + \sum_{i=2}^p \frac{1}{\log_2 i}$$

Finally, normalized discounted cumulative gain at rank p ($NDCG_p$) is expressed as:

$$NDCG_p = \frac{DCG_p}{IDCG_p}$$

with rel_i being the relevancy for document i in the retrieval ranking. For this experiment, this value is either 0 or 1.

1.2.5 Calculating Reciprocal Rank

Reciprocal rank is defined as the reciprocal of the rank at which the first relevant document is found, so if the 3^{rd} document in the retrieval ranking list is the first relevant document, the reciprocal rank is $\frac{1}{3}$.

1.3 Results

After building the index, CACM query 10 was processed by the `getrel.py` script, the output of which can be found in Listing 1. This script calculates all the values shown in Table 1, which are all of the required values for the question.

```
1 [mchaney@mchaney-l getrel]$ python q83.py -q 10 -n 10000
2 query 10
3 query: parallel languages languages for parallel computation
4 precision: 0.0190816935003
5 recall: 0.914285714286
6 precision @10: 0.9
7 NDCG @5: 1.0
8 NDCG @10: 0.942709999032
9 avg precision: 0.5922383982
10 reciprocal rank: 1.0
```

Listing 1: Output from running the `getrel.py` script for queries 1 and 10 from the CACM collection.

Query #	Avg. Prec.	NDCG @5	NDCG @10	Prec. @10	Recip. Rank
10	0.5922383982	1.0	0.942709999032	0.9	1.0

Table 1: Calculations for CACM query 10 from all retrieved documents.

2 Question 8.4

2.1 Question

For two queries in the CACM collection, generate two uninterpolated recall-precision graphs, a table of interpolated precision values at standard recall levels, and the average interpolated recall-precision graph.

2.2 Approach

Using the `getrel.py`, `q84.py` and `graphs.R` scripts, found in Listings 3, 5 and 8 were created to complete this task.

2.3 Results

2.3.1 Uninterpolated Recall-Precision Graph

The uninterpolated recall-precision graph is shown in Figure 1.




Figure 1: Uninterpolated Recall-Precision Graph for CACM Queries 6 and 8.

2.3.2 Interpolated Precision

The graph for the interpolated precision at standard recall values is shown in Figure 2 and the table of the values for each query, including the averages, is shown in Table 2.




Figure 2: Graph of interpolated precision at standard recall values for CACM queries 6 and 8.

Recall	0.0	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1.0
Query 6	0.333	0.333	0.333	0.333	0.333	0.333	0.333	0.333	0.333	0.333	0.333
Query 8	0.5	0.5	0.5	0.5	0.4	0.4	0.4	0.4	0.4	0.4	0.4
Average	0.417	0.417	0.417	0.417	0.367	0.367	0.367	0.367	0.367	0.367	0.367

Table 2: Interpolated precision at standard recall values for CACM queries 6 and 8.

2.3.3 Average Interpolated Precision

The graph of the average interpolated precision at standard recall values for CACM queries 6 and 8 can be found in Figure 3.




Figure 3: Average interpolated recall-precision graph for CACM Queries 6 and 8.

3 Question 8.5

3.1 Question

Generate the mean average precision, recall-precision graph, average NDCG at 5 and 10, and precision at 10 for the entire CACM query set.

3.2 Approach

The `getrel.py` and `q85.py` scripts, found in Listings 3 and 6, were used to complete this question.

3.3 Results

The output from running the `q85.py` script can be found in Listing ??.

Using only queries for which relevance judgments exist MAP, NDCG @5 and 10, and the precision @ 10 were calculated. The results can be found in Table 3. The generated recall-precision graph for the entire query set can be found in Figure 4.

MAP	NDCG @5	NDCG @10	Prec. @10
0.339552098123	0.461648777763	0.381724764912	0.317647058824

Table 3: Calculations for all CACM queries from all retrieved documents.




Figure 4: Recall-precision graph for all CACM Queries.

4 Question 8.7

4.1 Question

Another measure that has been used in a number of evaluations is R-precision. This is defined as the precision at R documents, where R is the number of relevant documents for a query. It is used in situations where there is a large variation in the number of relevant documents per query. Calculate the average *R-precision* for the CACM query set and compare it to the other measures.

4.2 Approach

The `getrel.py` and `q87.py` scripts were used to complete this exercise. They can be found in Listings 3 and 7. The script was run over the entire document set, which will minimize precision and maximize recall. This will be taken into consideration when comparing to the R-precision score.

4.3 Results

The output of running the `q87.py` script can be found in Listing 2.

```
1 [mchaney@mchaney-l getrel]$ python q87.py -n 3204 -q 10
2 query 10
3 query: parallel languages languages for parallel computation
4 relevant: 35
5 retrieved: 1677
6 precision: 0.0190816935003
7 recall: 0.914285714286
8 precision @10: 0.9
9 NDCG @5: 1.0
10 NDCG @10: 0.942709999032
11 avg precision: 0.5922383982
12 reciprocal rank: 1.0
13 |R|: 35
14 R-precision: 0.555555555556
```

Listing 2: Output of `q87.py` for query 10.

4.3.1 Comparison

For query 10 the R-precision score was 0.55, which is very close to the average precision over the entire document set. This rather simplistic comparison is one mark towards R-precision being a decently reliable measure of performance for a ranking. R-precision seems to be somewhat of an intuitive normalization of the precision. This is because the sample size of documents is bounded by the size of the relevant set of documents. This measure will favor rankings that push more relevant documents into higher ranks, which seems like a strong measure.

The only problem with this measure is if the relevant set is very small the results of the calculation may vary wildly. This may make the measure inappropriate for a targeted search because it will basically be 0 or 1, which isn't useful for comparing rankings.

5 Appendix

5.1 Code listings

```
1 import argparse
2 import re
3 import requests
4 import sys
5 import xmltodict
6 import numpy as np
7 from math import log
8 from bs4 import BeautifulSoup
9
10 def parseargs():
11     parser = argparse.ArgumentParser()
12     parser.add_argument('-p', '--port', type=int, default=42247, help='galago server port')
13     parser.add_argument('-q', '--qnum', nargs='+', type=int, default=[6, 8], help='query
        number')
14     parser.add_argument('-n', type=int, default=10, help='number of retrieval pages')
15     return parser.parse_args()
16
17 args = parseargs()
18
19 def buildrel():
20     rel = {}
21     for line in open('cacm.rel').readlines():
22         q, _, doc, _ = line.split()
23         if q not in rel:
24             rel[q] = []
25         rel[q].append(int(doc.split('-')[1]))
26     return rel
27
28 def buildqueries():
29     with open('cacm.query.xml') as fd:
30         return xmltodict.parse(fd.read())
31
32 REL = buildrel()
33 QUERIES = buildqueries()
34 RE = re.compile('/home/mchaney/workspace/edu/cs834-f16/assignments/assignment4/code/cacm/
        docs/CACM-([d]+).html')
35 ID = {'id': 'result'}
36 URL = 'http://0.0.0.0:{0}/search'
37 QUERY1 = 'what articles exist which deal with tss time sharing system an operating system
        for ibm computers'
38 PDICT = {'q': QUERY1, 'start': 0, 'n': args.n}
39
40 def query(qstr, port=args.port):
41     PDICT['q'] = qstr
42     PDICT['n'] = args.n
43     res = requests.get(URL.format(port), params=PDICT)
44     if not res.ok:
45         return None
46     soup = BeautifulSoup(res.text, 'html.parser')
47     return [int(RE.match(href.text).groups()[0]) for href in soup.select("#result a")]
48
49 def recall(rel, retr):
50     relset = set(rel)
51     retrset = set(retr)
52     return float(len(relset.intersection(retrset))) / len(relset)
53
54 def precision(rel, retr):
55     relset = set(rel)
56     retrset = set(retr)
57     return float(len(relset.intersection(retrset))) / len(retrset)
58
59 def run(rel, retr, func):
60     rr = []
61     for i in range(1, len(retr)+1):
62         rr.append(func(rel, retr[:i]))
63     return rr
64
65 def avg(rel, retr, func):
66     prun = run(rel, retr, func)
67     res = []
68     for i in range(len(retr)):
```

```

69         if retr[i] in rel:
70             res.append(prun[i])
71     if len(res) == 0:
72         return 0.0
73     return float(sum(res))/len(res)
74
75 def getrel(rel, retr, i):
76     return 1 if retr[i] in rel else 0
77
78 def DCG(rel, retr, p):
79     sum = 0
80     for i in range(2, p+1):
81         sum += float(getrel(rel, retr, i-1)) / log(i, 2)
82     return getrel(rel, retr, 0) + sum
83
84 def IDCG(p):
85     sum = 0
86     for i in range(2, p+1):
87         sum += 1 / log(i, 2)
88     return 1 + sum
89
90 def NDCG(rel, retr, p):
91     dcg = DCG(rel, retr, p)
92     idcg = IDCG(p)
93     return dcg / idcg
94
95 def reciprank(rel, retr):
96     for i in range(1, len(retr)+1):
97         if retr[i-1] in rel:
98             return 1.0 / i
99     return 0.0
100
101 def ipr(rrun, prun):
102     res = []
103     for i in np.arange(0, 1.1, .1):
104         for j in range(len(rrun)):
105             if rrun[j] > i:
106                 idx = j
107                 break
108     res.append(max(prun[idx:]))
109     return np.arange(0, 1.1, 0.1), res
110
111 def rprecision(rel, prun):
112     return prun[len(rel)]
113
114 def getquery(qnum):
115     return QUERIES['parameters'][ 'query' ][qnum-1][ 'text' ]
116
117 def process(qnum):
118     qstr = getquery(qnum)
119     retr = query(qstr)
120     if str(qnum) not in REL:
121         return [None]*12
122     rel = REL[str(qnum)]
123     prun = run(rel, retr, precision)
124     rrun = run(rel, retr, recall)
125     prec = precision(rel, retr)
126     rec = recall(rel, retr)
127     avgprec = avg(rel, retr, precision)
128     ndcg5 = NDCG(rel, retr, 5)
129     ndcg10 = NDCG(rel, retr, 10)
130     recip = reciprank(rel, retr)
131     return qnum, qstr, retr, rel, prun, rrun, prec, rec, ndcg5, ndcg10, avgprec, recip
132
133 def printresults(qnum, qstr, retr, rel, prun, rrun, prec, rec, ndcg5, ndcg10, avgprec, recip):
134     if not qnum:
135         return
136     print 'query {0}'.format(qnum)
137     print 'query: {0}'.format(qstr)
138     if args.n == 10:
139         print 'relevant: {0}'.format(len(rel))
140         print 'retrieved: {0}'.format(len(retr))
141         print 'p-run: {0}'.format(prun)
142         print 'r-run: {0}'.format(rrun)
143     else:
144         print 'relevant: {0}'.format(len(rel))

```

```

145     print 'retrieved: {}'.format(len(retr))
146     print 'precision: {}'.format(prec)
147     print 'recall: {}'.format(rec)
148     print 'precision @10: {}'.format(prun[9])
149     print 'NDCG @5: {}'.format(ndcg5)
150     print 'NDCG @10: {}'.format(ndcg10)
151     print 'avg precision: {}'.format(avgprec)
152     print 'reciprocal rank: {}'.format( recip)
153
154 def printdata(rrun, prun, fname):
155     with open(fname, 'w') as fd:
156         zipped = zip(rrun, prun)
157         for z in zipped:
158             fd.write('{}\t{}\n'.format(z[0], z[1]))
159
160 if __name__ == '__main__':
161     for qnum in args.qnum:
162         printresults(*process(qnum))

```

Listing 3: getrel.py

```

1 from getrel import *
2
3 TABLE = """\begin{table}[h!]
4 \\\centering
5 \\\begin{tabular}{|c|c|c|c|c|c|c|c|}
6 \\\hline
7 Query \# & Avg. Prec. & NDCG @5 & NDCG @10 & Prec. @10 & Recip. Rank \\\
8 \\\hline
9 {0} & {1} & {2} & {3} & {4} & {5} \\\
10 \\\hline
11 \\\end{tabular}}
12 \\\caption{{Calculations for CACM query {6} from all retrieved documents.}}
13 \\\label{{tab:q83}}
14 \\\end{table}}
15 """
16
17 def printtab(qnum, qstr, retr, rel, prun, rrun, prec, rec, ndcg5, ndcg10, avgprec, recip,
18             rprec):
19     fname = 'query{0}.tab'.format(qnum)
20     with open(fname, 'w') as fd:
21         fd.write(TABLE.format(qnum, avgprec, ndcg5, ndcg10, prun[9], recip, qnum))
22
23 for qnum in args.qnum:
24     results = process(qnum)
25     printresults(*results)
26     printtab(*results)

```

Listing 4: q83.py

```

1 from getrel import *
2
3 HEAD = """\begin{table}[H]
4 \\\centering
5 \\\begin{tabular}{l l l l l l l l l l l l}
6 Recall & 0.0 & 0.1 & 0.2 & 0.3 & 0.4 & 0.5 & 0.6 & 0.7 & 0.8 & 0.9 & 1.0 \\\
7 \\\cline{2-12}
8 """
9
10 ROW = """{0} & {1:.3g} & {2:.3g} & {3:.3g} & {4:.3g} & {5:.3g} & {6:.3g} & {7:.3g} & {8:.3g}
11 & {9:.3g} & {10:.3g} & {11:.3g} \\\
12 \\\cline{{2-12}}
13 """
14
15 TAIL = """\end{tabular}
16 \\\caption{{.}}
17 \\\label{{tab:ipr68}}
18 \\\end{table}}
19 """
20
21 def printtable(iprl):
22     with open('iptab.tex', 'w') as fd:
23         fd.write(HEAD)
24         for iprun, qnum in iprl:

```

```

24         fd.write(ROW.format('Query {0}'.format(qnum), *iprun))
25         avg = [float(sum(col))/len(col) for col in zip(*[col[0] for col in iprl])]
26         printdata(np.arange(0, 1.1, .1), avg, 'avg.dat')
27         fd.write(ROW.format('Average', *avg))
28         fd.write(TAIL)
29
30     iprl = []
31     for qnum in args.qnum:
32         results = process(qnum)
33         printresults(*results)
34         qnum, qstr, retr, rel, prun, rrun, prec, rec, ndcg5, ndcg10, avgprec, recip, rprec =
            results
35         printdata(rrun, prun, 'urpg{0}.dat'.format(qnum))
36         irrun, iprun = ipr(rrun, prun)
37         printdata(irrun, iprun, 'ipr{0}.dat'.format(qnum))
38         iprl.append((iprun, qnum))
39     printtable(iprl)

```

Listing 5: q84.py

```

1  from getrel import *
2
3  TABLE = """\begin{table}[h!]
4  \centering
5  \begin{tabular}{|c|c|c|c|}
6  \hline
7  MAP & NDCG @5 & NDCG @10 & Prec. @10 \\
8  \hline
9  {0} & {1} & {2} & {3} \\
10 \hline
11 \end{tabular}
12 \caption{{Calculations for all CACM queries from all retrieved documents.}}
13 \label{{tab:q85}}
14 \end{table}
15 """
16
17 def printtab(fname, cacmmap, avgndcg5, avgndcg10, avgprec10):
18     with open(fname, 'w') as fd:
19         fd.write(TABLE.format(cacmmap, avgndcg5, avgndcg10, avgprec10))
20
21     netavg = []
22     iprl = []
23     ndcg5lst = []
24     ndcg10lst = []
25     prunlst = []
26     for i in range(1, 64):
27         qnum, qstr, retr, rel, prun, rrun, prec, rec, ndcg5, ndcg10, avgprec, recip = process(i)
28         if avgprec:
29             netavg.append(avgprec)
30             prunlst.append(prun[9])
31             irrun, iprun = ipr(rrun, prun)
32             iprl.append((iprun, qnum))
33             ndcg5lst.append(ndcg5)
34             ndcg10lst.append(ndcg10)
35
36     # MAP
37     cacmmap = float(sum(netavg)) / len(netavg)
38     print 'average precision: {0}'.format(cacmmap)
39
40     # Recall-Precision
41     netavgrpg = [float(sum(col))/len(col) for col in zip(*[col[0] for col in iprl])]
42     printdata(np.arange(0, 1.1, .1), netavgrpg, 'avgq85.dat')
43
44     # NDCG @ 5 and 10
45     avgndcg5 = float(sum(ndcg5lst))/len(ndcg5lst)
46     avgndcg10 = float(sum(ndcg10lst))/len(ndcg10lst)
47     print 'NDCG @5 : {0}'.format(avgndcg5)
48     print 'NDCG @10: {0}'.format(avgndcg10)
49
50     # precision at 10
51     avgprec10 = float(sum(prunlst))/len(prunlst)
52     print 'Precision @10: {0}'.format(avgprec10)
53
54     printtab('q85.tab', cacmmap, avgndcg5, avgndcg10, avgprec10)

```

Listing 6: q85.py

```

1 from getrel import *
2
3 for qnum in args.qnum:
4     qnum, qstr, retr, rel, prun, rrun, prec, rec, ndcg5, ndcg10, avgprec, recip = process(
5         qnum)
6     printresults(qnum, qstr, retr, rel, prun, rrun, prec, rec, ndcg5, ndcg10, avgprec, recip)
7     rprec = rprecision(rel, prun)
8     print '|R|: {}'.format(len(rel))
9     print 'R-precision: {}'.format(rprec)

```

Listing 7: q87.py

```

1 plotone <- function(data, fname) {
2     pdf(fname)
3     plot(data, type='o', pch=15, ylim=c(0,1), xlim=c(0,1),
4         ylab="Precision", xlab="Recall")
5     dev.off()
6 }
7 urpgraph <- function(d1, d2, fname) {
8     pdf(fname)
9     plot(d1, lwd=2, type='o', pch=18, ylim=c(0,1), xlim=c(0,1), col="gray60",
10        ylab="Precision", xlab="Recall")
11     lines(d2, lwd=2, type="o", pch=15, col="gray30")
12     legend(0.8, 1, c('Query 6', 'Query 8'), cex=0.8,
13        col=c('gray60', 'gray30'), lty=c(1,1), pch=c(18,15))
14     dev.off()
15 }
16 iprgraph <- function(d1, d2, id1, id2, fname) {
17     pdf(fname)
18     plot(d1, lwd=2, type="p", pch=18, ylim=c(0,1), xlim=c(0,1), col="gray60",
19        ylab="Precision", xlab="Recall")
20     lines(id1, lwd=2, type="s", col="gray60")
21     lines(d2, lwd=2, type="p", pch=15, col="gray30")
22     lines(id2, lwd=2, type="s", col="gray30")
23     legend(0.8, 1, c('Query 6', 'Query 8'), cex=0.8,
24        col=c('gray60', 'gray30'), lty=c(1,1), pch=c(18,15))
25     dev.off()
26 }
27 aipgraph <- function(avg, id1, id2, fname) {
28     pdf(fname)
29     plot(avg, lwd=2, type="l", ylim=c(0,1), xlim=c(0,1), col="black",
30        ylab="Precision", xlab="Recall")
31     lines(id1, lwd=2, type="s", col="gray60")
32     lines(id2, lwd=2, type="s", col="gray30")
33     lines(id2, lwd=2, type="s", col="gray30")
34     legend(0.8, 1, c('Query 6', 'Query 8', 'Average'), cex=0.8,
35        col=c('gray60', 'gray30', 'black'), lty=c(1,1,1), pch=c(18,15,17))
36     dev.off()
37 }
38
39 args = commandArgs(trailingOnly=TRUE)
40
41 d1 <- read.table(paste('urpg', args[1], '.dat', sep=''))
42 d2 <- read.table(paste('urpg', args[2], '.dat', sep=''))
43
44 plotone(d1, paste('urpg', args[1], '.pdf', sep=''))
45 plotone(d2, paste('urpg', args[2], '.pdf', sep=''))
46 urpgraph(d1, d2, paste('urpg', args[1], '', args[2], '.pdf', sep=''))
47
48 id1 <- read.table(paste('ipr', args[1], '.dat', sep=''))
49 id2 <- read.table(paste('ipr', args[2], '.dat', sep=''))
50 iprgraph(d1, d2, id1, id2, paste('ipr', args[1], '', args[2], '.pdf', sep=''))
51
52 avg <- read.table('avg.dat')
53 aipgraph(avg, id1, id2, paste('aipr', args[1], args[2], '.pdf', sep=''))
54
55 overallavg <- read.table('avgq85.dat')
56 plotone(overallavg, 'avgq85.pdf')

```

Listing 8: Script used to generate the recall-precision graphs

6 References

- [1] Kenneth Reitz. Requests: HTTP for Humans. Available at <http://docs.python-requests.org/en/master/>. Accessed: 2016/09/20.
- [2] Leonard Richardson. Beautiful Soup. Available at: <https://www.crummy.com/software/beautifulsoup/>. Accessed: 2016/09/20.