Cluster-Based Retrieval Using Language Models

Proceedings of the 27th Annual International ACM SIGIR Conference on Research and Development in Information Retrieval. ACM, 2004.

Xiaoyong Liu and W. Bruce Croft.

A Cluster-Based Resampling Method for Pseudo-Relevance Feedback

Proceedings of the 31st Annual International ACM SIGIR Conference on Research and Development in Information Retrieval. ACM, 2008.

Kyung Soon Lee, W. Bruce Croft, and James Allan

Presented by Matt Chaney

CS 834 - Presentation 4

Motivation

- Similar documents can serve the same information need
- Many studies applying clustering to IR
  - Initially to improve effectiveness/efficiency or categorize documents
- Previous study lacked definition for finding optimal document clusters
  - Automatically
  - Without relevance judgements
  - On collections of realistic size
- Apply new language modeling techniques to cluster-based IR systems
  - Provide principled way to explore document-cluster relationships
  - Language models allow for use of sophisticated smoothing parameters

Clustering

- Grouping things based on similarity of features

- Dogs
  - Height
  - Weight

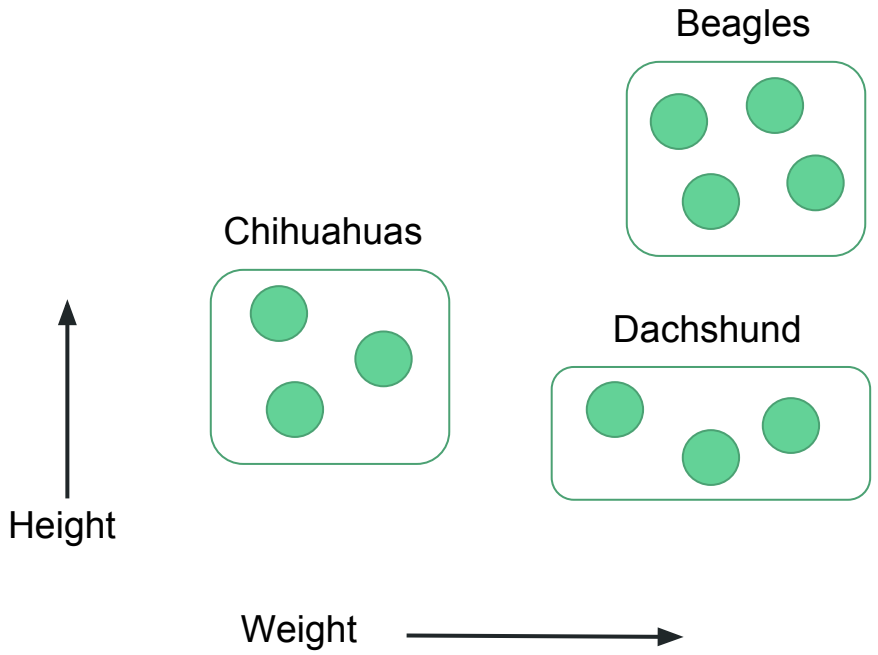

Language Models

Probability distribution over all terms in a language vocabulary

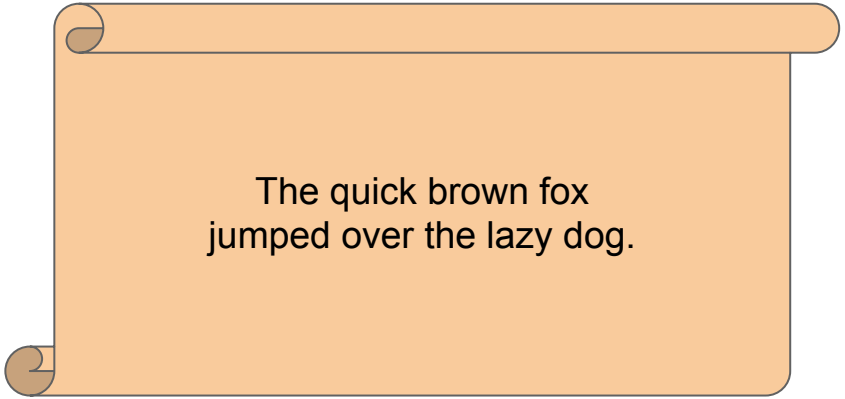

Word	Probability
The	0.222
Quick	0.111
Brown	0.111
Fox	0.111
⋮	⋮

Cluster-based Language Models

- Organize document collections around topics
- Create language model for topics to use as representation
- Estimate query-likelihood using cluster topics and select collection with best topic

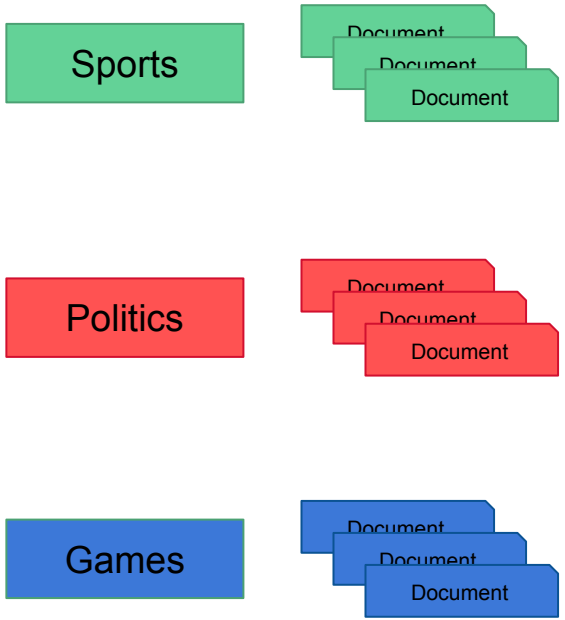

Cluster-based information retrieval

- Rank clusters in response to query (CQL)
  - By calculating centroid of cluster in comparison to query terms
  - Assumed arbitrary document in higher ranked cluster more relevant

- Use clusters as a form of *document smoothing* (CBDM)
  - Grouping similar documents smooths out differences among individuals

Traditional Query-Likelihood (QL)

- Standard document language model

$$P(Q \mid D) = \prod_{i=1}^{m} P(q_i \mid D)$$

- Adding smoothing

$$P(w \mid D) = \lambda P_{ML}(w \mid D) + (1 - \lambda) P_{ML}(w \mid Coll)$$

- Simple Jelinek-Mercer

- Bayesian with Dirichlet prior

$$\lambda = \frac{\sum_{w' \in D} tf(w', D)}{\sum_{w' \in D} tf(w', D) + \mu}$$

Cluster-Based Query Likelihood (CQL)

- Calculate probability of query given a cluster language model

$$P(Q \mid Cluster) = \prod_{i=1}^{m} P(q_i \mid Cluster)$$

- Cluster language model

$$P(w \mid Cluster) = \lambda P_{ML}(w \mid Cluster) + (1-\lambda) P_{ML}(w \mid Coll)$$

$$= \lambda \frac{tf(w, Cluster)}{\sum_{w' \in cluster} tf(w', Cluster)} + (1-\lambda) \frac{tf(w, Coll)}{\sum_{w' \in V} tf(w', Coll)}$$

Cluster-Based Document Smoothing (CBDM)

- Smooth document language model based on similar documents

$$P(w|D) = \lambda P_{ML}(w|D) + (1-\lambda)P(w|Cluster)$$
$$= \lambda P_{ML}(w|D) + (1-\lambda)[\beta P_{ML}(w|Cluster) + (1-\beta)P_{ML}(w|Coll)]$$

- Both λ and β are general symbols for smoothing

- Two-stage smoothing
  - Cluster model smoothed with collection model
  - Document model smoothed with smoothed cluster model

Clustering Approaches

- Distance Measures
  - Dice, Jaccard and overlap coefficients
  - Kullback-Liebler (KL) Divergence
  - **Cosine Measure**
- Partitioning ➜ Static Clustering
  - **Three-pass K-means clustering**
- Hierarchical Agglomerative ➜ Query-specific Clustering
  - Single/Complete Linkage
  - Group Average
  - Centroid
  - Ward’s method

Experimental Data

All queries taken from *title* field of TREC topics

Collection	Contents	# of Docs	Size	Average # of Words/Doc[1]	Queries	# of Queries with Relevant Docs
AP	Associated Press newswire 1988-90	242,918	0.73 Gb	473.6	TREC topics 51-150 (title only)	99
FR	Federal Register 1988-89	45,820	0.47 Gb	873.9	TREC topics 51-100 (title only)	21
WSJ	Wall Street Journal 1987-92	173,252	0.51 Gb	465.8	TREC topics 51-100 & 151-200 (title only)	100
FT	Financial Times 1991-94	210,158	0.56 Gb	412.7	TREC topics 301-400 (title only)	95
SJMN	San Jose Mercury News 1991	90,257	0.29 Gb	453.0	TREC topics 51-150 (title only)	94
LA	LA Times	131,896	0.48 Gb	526.5	TREC topics 301-400 (title only)	98

Table 1. Statistics of data sets.
Cluster-Based Retrieval Using Language Models. Xiaoyong Liu and W. Bruce Croft. 2004.

Experimental Design

- Evaluate ranking clusters (CQL method) with the AP and WSJ collections
  - Five clustering algorithms for cluster language model compared to baseline (QL+DM)
  - Bayesian smoothing w/ Dirichlet prior
  - Jelinek-Mercer smoothing
- Cluster-based retrieval (CBDM)
  - Query-Likelihood (QL+DM)
  - Relevance Model (RM+DM)
- Measured in Average Precision
- Trained parameter values before actual tests
  - Various settings for cluster distance threshold as well as smoothing parameters

Cluster-based IR by Ranking Clusters (CQL)

1. Document-based, query-likelihood retrieval
2. Cluster top 1,000 results
3. Rank clusters with CQL method
4. Return ordered list of clusters where documents within cluster ranked according to step 1.

Collection	First-stage doc retrieval (QL+DM)	Group-average	Single-linkage	Complete-linkage	Centroid	Ward's
AP (training)	0.2179	0.2161 (t=0.8)	0.2153 (t=0.8)	0.2130 (t=0.8)	0.2164 (t=0.7)	0.2160 (t=0.8)
WSJ	0.2958	0.2902 (t=0.8)	0.2911 (t=0.8)	0.2889 (t=0.8)	0.2936 (t=0.8)	0.2963 (t=0.8)

Table 2. CQL results.
Cluster-Based Retrieval Using Language Models. Xiaoyong Liu and W. Bruce Croft. 2004.

Cluster-based Document Smoothing (CBDM)

Collection	Simple Okapi	QL+DM	QL+CBDM	%chg	RM+DM	RM+CBDM	%chg
AP (K=2000)	0.2198	0.2179	0.2326 (+)	+6.73*	0.2745	0.2775	+1.08
WSJ (K=2000)	0.2762	0.2958 (+)	0.3006 (+)	+1.62*	0.3422	0.3445	+0.64
FT (K=2000)	0.2556	0.2610	0.2713 (+)	+3.95*	0.2835	0.2845	+0.36
SJMN (K=2000)	0.2098	0.2032	0.2171 (+)	+6.88*	0.2633	0.2673	+1.52*
LA (K=2000)	0.2279	0.2468 (+)	0.2590 (+)	+4.94*	0.2614	0.2621	+0.28
FR (K=1000)	0.2644	0.2875	0.3316	+15.37	0.1486	0.1934	+30.10

Table 5. Evaluation of Cluster-based Retrieval compared with Simple Okapi method[1]
Cluster-Based Retrieval Using Language Models. Xiaoyong Liu and W. Bruce Croft. 2004.

[1] Sparck et. al. A probabilistic model of information retrieval: development and comparative experiments. 2004.

Cluster-Based Retrieval Using Language Models

Proceedings of the 27th Annual International ACM SIGIR Conference on Research and Development in Information Retrieval. ACM, 2004.

Xiaoyong Liu and W. Bruce Croft.

A Cluster-Based Resampling Method for Pseudo-Relevance Feedback

Proceedings of the 31st Annual International ACM SIGIR Conference on Research and Development in Information Retrieval. ACM, 2008.

Kyung Soon Lee, W. Bruce Croft, and James Allan

Presented by Matt Chaney

CS 834 - Presentation 4

Pseudo-Relevance Feedback (PRF)

- PRF assumes top-retrieved documents are relevant
  - Query expansion
  - Re-evaluate initial rankings

- Top retrieved documents can contain “noise” documents
  - Sub-optimal precision allows non-relevant documents to sway relevancy judgment away from initial query

- Two central issues to utilizing PRF to improve query effectiveness
  - How to select relevant documents from initial retrieval set
  - How to select query expansion terms

Recent Cluster-based IR and PRF Developments

- Cluster theory applied successfully to IR
  - Re-ranking using clusters
  - **Cluster-based retrieval**
  - Score regularization

- Sampling / Resampling - Select more varied and novel top-ranked docs for PRF
  - Random
  - **Selective**
    - Skipping some top-ranked documents
    - Boosting - Focus subsequent training on poor performers
    - Query variant result resampling

Cluster-Based Selective Resampling

- Resampling method using clusters
  - Document clusters can represent query subtopics
  - *Dominant documents* appear in overlapping clusters
  - Selectively resampling these documents emphasizes core query topics

- Cluster-based resampling can achieve a higher *relevance density*

$$Density = \frac{the\ \ number\ \ \ of\ \ relevant\ \ feedback\ \ documents}{the\ \ number\ \ of\ \ feedback\ \ \ documents}$$

Resampling Process

- Uses language model and relevance model frameworks
- Dominant documents contribute more to expansion terms than other documents

Process:

1. Retrieve initial results using query-likelihood language model

$$P(Q \mid D) = \prod_{i=1}^{m} P(q_i \mid D) \longrightarrow$$

$$P(w \mid D) = \frac{|D|}{|D| + \mu} P_{ML}(w \mid D) + \frac{\mu}{|D| + \mu} P_{ML}(w \mid Coll)$$

$$P_{ML}(w \mid D) = \frac{freq(w, D)}{|D|}, \quad P_{ML}(w \mid Coll) = \frac{freq(w, Coll)}{|Coll|}$$

Clustering for Dominant Documents

2. Cluster these results using K-nearest neighbors to find dominant documents
    - Assumption: If a document is in several clusters that are highly related to the query it is considered a *dominant document*

To which group do we assign the green point?

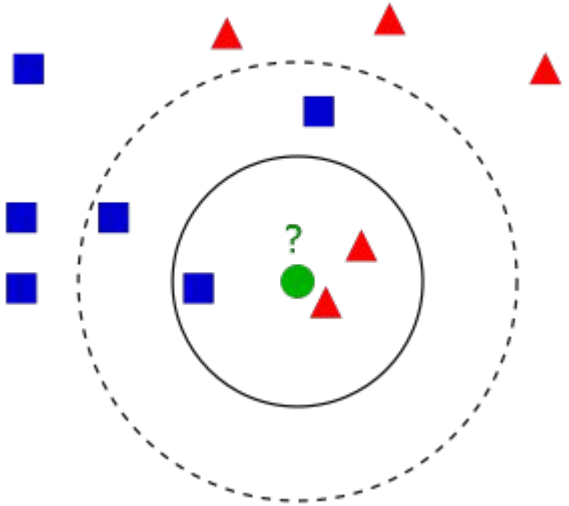

Fig 1. Example of k-NN classification.
Creative Commons. Public Domain.

Using cluster-based ranking

3. Rank clusters using **cluster-based ranking**
   - Treating a document cluster as a large document allows use of the query-likelihood model

$$P(Q \mid Clu) = \prod_{i=1}^{m} P(q_i \mid Clu)$$

$$P(w \mid Clu) = \frac{|Clu|}{|Clu| + \lambda} P_{ML}(w \mid Clu) + \frac{\lambda}{|Clu| + \lambda} P_{ML}(w \mid Coll)$$

$$P_{ML}(w \mid Clu) = \frac{freq(w, Clu)}{|Clu|}, \quad P_{ML}(w \mid Coll) = \frac{freq(w, Coll)}{|Coll|}$$

Query Expansion Term Selection

4. Select query terms from each document in top-ranked clusters using the relevance model

$$\sum_{D \in R} P(D)P(w \mid D)P(Q \mid D)$$

Experimental Setup

- Several Test Collections
- Large heterogeneous web collections
  - GOV2
  - WT10G
- Small homogeneous news-related collections
  - AP
  - WSJ
  - ROBUST
- Topic title field used as query

Collection Summary

Collection	Description	# of docs	Topics	
			Train	Test
GOV2	2004 crawl of .gov domain	25,205,179	701-750	751-800
WT10g	TREC web collection	1,692,096	451-500	501-550
ROBUST	Robust 2004 collection	528,155	301-450	601-700
AP	Association Press 88-90	242,918	51-150	151-200
WSJ	Wall street Journal 87-92	173,252	51-150	151-200

Table 1. Summary of Test Collections.
A Cluster-Based Resampling Method for Pseudo-Relevance Feedback. Lee, Croft, Allan. 2008.

Training

- Each collection is divided into training and testing topics

- Training used to tune parameters
  - μ in initial query-likelihood model $\quad \mu \in \{ 500, 750, 1000, 1500, 2000, \ldots, 5000 \}$
  - Number of feedback documents $\quad |R| \in \{ 5, 10, 25, 50, 75, 100 \}$
  - Number of expansion terms $\quad e \in \{ 10, 25, 50, 75, 100 \}$
  - Weight of original query $\quad \lambda \in \{ 0.1, 0.2, \ldots, 0.9 \}$
  - Number of clusters $\quad |C| \in \{ 1, 2, 5, 10, 15, 20 \}$

- Combined expansion terms with query via *Indri form*

$$\#weight\ (\lambda\ \#combine\ (q_1 \ldots q_m)$$
$$(1-\lambda)\#weight\ (p_1\ t_1 \ldots p_e\ t_e))$$

Experimental Comparisons

- Baseline models
  - Language Model (LM)
  - Relevance Model (RM)
- Cluster-based Reranking Method (Rerank)
- **Cluster-based Resampling**
- Upper Bound - True relevance feedback (TrueRF)

Test Collection Results

	LM	Rerank	RM	Resampling	TrueRF
GOV2	0.3258	0.3406 $^{\alpha}$	0.3581 $^{\alpha\beta}$	0.3806 $^{\alpha\beta\gamma}$	0.4315 $^{\alpha\beta\gamma\delta}$
WT10g	0.1861	0.2044 $^{\alpha}$	0.1966	0.2352 $^{\alpha\beta\gamma}$	0.4030 $^{\alpha\beta\gamma\delta}$
ROBUST	0.2920	0.3206 $^{\alpha}$	0.3591 $^{\alpha\beta}$	0.3515 $^{\alpha\beta}$	0.5351 $^{\alpha\beta\gamma\delta}$
AP	0.2077	0.2361 $^{\alpha}$	0.2803 $^{\alpha\beta}$	0.2906 $^{\alpha\beta}$	0.4253 $^{\alpha\beta\gamma\delta}$
WSJ	0.3258	0.3611 $^{\alpha}$	0.3967 $^{\alpha\beta}$	0.4033 $^{\alpha\beta}$	0.5306 $^{\alpha\beta\gamma\delta}$

Table 2. Performance comparisons using MAP for test topics on test collections..
A Cluster-Based Resampling Method for Pseudo-Relevance Feedback. Lee, Croft, Allan. 2008.

Relevance Density

- Explain why this method works
- Compare Cluster-based resampling to PRF without redundant document resampling
- Recall

$$Density = \frac{the\ \ number\ \ \ of\ \ \ relevant\ \ \ feedback\ \ documents}{the\ \ number\ \ of\ \ feedback\ \ \ documents}$$

Relevance Density Measure

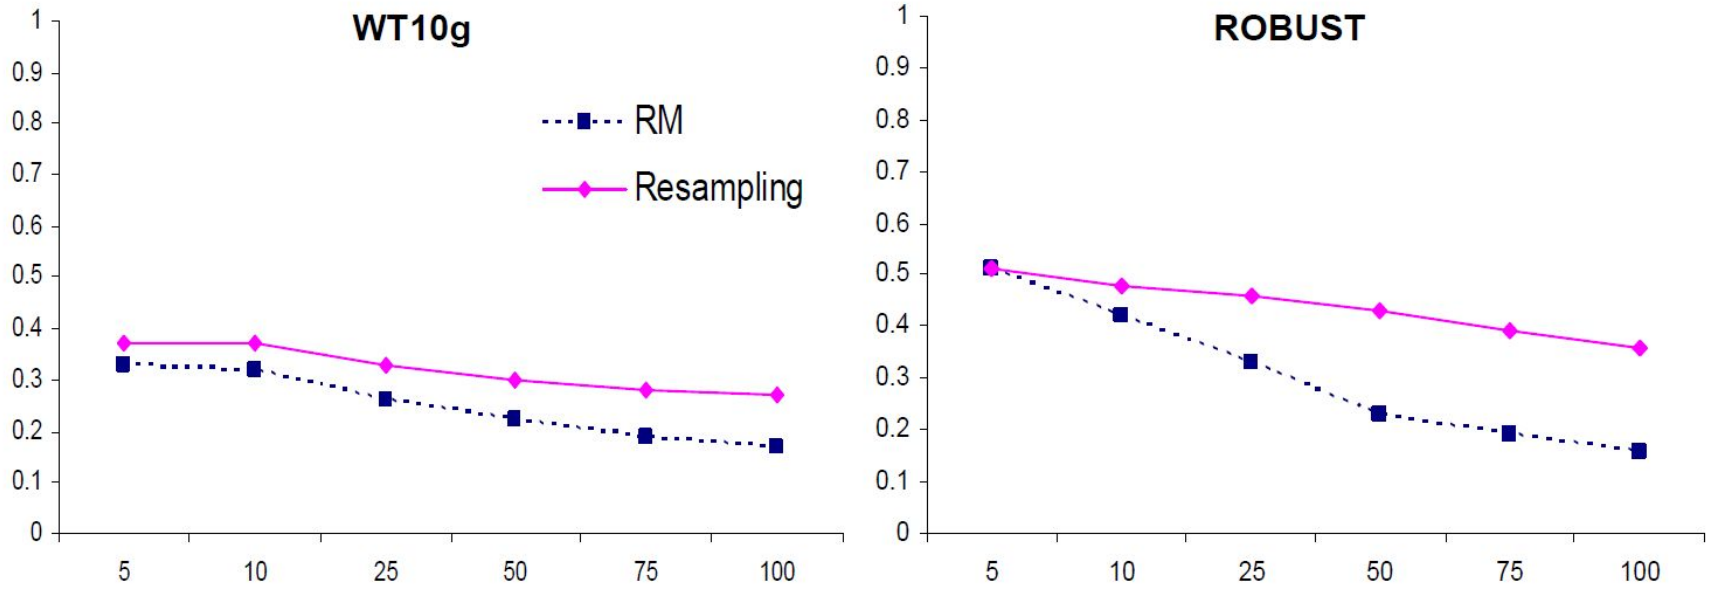

Fig 1. The relevance density for RM and Resampling
A Cluster-Based Resampling Method for Pseudo-Relevance Feedback. Lee, Croft, Allan. 2008.

Relevance Density Results

	LM	RM	chg%	Resampling	chg%
GOV2	0.3258	0.3519 $^{\alpha}$	8.01	0.3764 $^{\alpha\beta}$	15.53
WT10G	0.1861	0.1886	1.34	0.2072 $^{\alpha}$	11.34
ROBUST	0.2920	0.3262 $^{\alpha}$	11.71	0.3549 $^{\alpha\beta}$	21.54
AP	0.2077	0.2758 $^{\alpha}$	32.79	0.2853 $^{\alpha}$	37.36
WSJ	0.3258	0.3785 $^{\alpha}$	16.18	0.4009 $^{\alpha\beta}$	23.05

Table 3. Performance on fixed feedback (set to 100 documents).
A Cluster-Based Resampling Method for Pseudo-Relevance Feedback. Lee, Croft, Allan. 2008.